

ChainIO: Seamless Kernel-Bypass for Composite Disk+Network Workloads via eBPF-Powered Syscall Chaining

Anonymous Authors

Abstract

Modern data-driven services like distributed analytical databases incur high tail-latencies because each storage operation (`read`) and network operation (`send/recv`) triggers a separate user-kernel crossing. While `io_uring` accelerates block I/O and `AF_XDP` accelerates packet I/O, no solution chains them end-to-end in a unified framework. We introduce ChainIO, a hybrid I/O-bypass framework that transparently live-patches POSIX I/O calls into a unified submission queue, and coordinates disk and network operations via shared BPF maps. By chaining these bypass paths with in-kernel eBPF programs, ChainIO achieves near zero-copy, batched I/O across domains, without modifying application source code. Our preliminary evaluation on ClickHouse TPC-H queries shows up to 4× higher IOPS and 30

1 Introduction

The performance cost of traditional system calls has become even more pronounced in the wake of Spectre and Meltdown mitigations, which add extra overhead to every user-kernel transition. This particularly impacts data-intensive applications like distributed analytical databases, where each operation may require multiple syscalls across both storage and networking domains.

A typical OLAP query in a distributed system like ClickHouse follows this path: client query → decomposition into column-file reads (`pread()`) → possible distribution to remote shards (`send()`) → network stack → remote node's `recv()` → context switch → disk lookup → response aggregation. Even on modern hardware, the aggregate query latency can reach hundreds of milliseconds due to these repeated context switches and data copies, particularly for fan-out queries that access dozens of partitions.

Recent kernel innovations have addressed individual domains: `io_uring` provides an asynchronous, batched interface for storage I/O, while `AF_XDP` enables zero-copy, kernel-bypass networking. However, these mechanisms remain siloed - `io_uring` still pays full cost for network operations, and `AF_XDP` doesn't accelerate disk access.

We introduce ChainIO, a unified syscall-chaining framework that fuses both bypass paths by:

- Dynamically rewriting POSIX I/O calls into batched `io_uring` submissions using bpftime-like uprobes.
- Bridging storage and network domains through shared BPF maps and zero-copy memory regions.
- Coordinating cross-domain operations to preserve correctness while minimizing context switches.
- Adaptively optimizing for tail latency through intelligent batching and prioritization.

ChainIO requires no application code changes - it transparently intercepts syscalls via binary rewriting and USDT probes, then redirects them through a unified bypass path. By chaining I/O operations across domains, ChainIO dramatically reduces query latency for composite workloads like distributed OLAP engines.

2 Background & Motivation

Modern storage and networking stacks offer domain-specific optimizations that fail to address composite workloads requiring both disk and network I/O. While `io_uring` provides asynchronous, batched disk operations, it still requires expensive syscalls for network traffic; conversely, `AF_XDP` delivers zero-copy network acceleration but offers nothing for storage access. Previous research has approached this challenge from several angles without fully solving the cross-domain problem: FlexSC [?] and MegaPipe batch syscalls but focus on single domains, while DPDK/SPDK and Demikernel [?] bypass the kernel entirely but require extensive application modifications. eBPF-based solutions like XRP [?] and BPF-oF [?] accelerate specific I/O paths (NVMe reads, remote storage) without addressing cross-domain dependencies. Even architectural innovations like IX [?] that redesign the OS with separate control and data planes remain siloed in network or storage specialization, leaving a critical gap for workloads that chain operations across both domains.

ClickHouse's `MergeTree` engine exemplifies this cross-domain problem, as shown in the profiling data in Figure ???. Its columnar storage engine issues large numbers of small, random `read()` calls against compressed column files and mark-file offsets, with each compressed-block fetch and metadata lookup translating into a user-kernel transition. In distributed setups, remote-shard




Figure 1: Flamegraph of ClickHouse Server showing syscall overhead

requests add further `send()` and `recv()` calls for data fetches and Raft heartbeats. Our profiling shows that the blocking `read()` syscall alone consumes 25% of query time, while small network receives (heartbeats, shard updates) account for another 2%. The cumulative cost of these syscalls—exacerbated by Spectre/Meltdown mitigations—introduces tens of microseconds of overhead per transition, multiplying into hundreds of milliseconds on fan-out queries. When a query spans dozens of remote partitions, each extra transition adds up quickly, creating a critical bottleneck for interactive dashboards and real-time analytics that cannot be solved by optimizing either storage or networking in isolation.

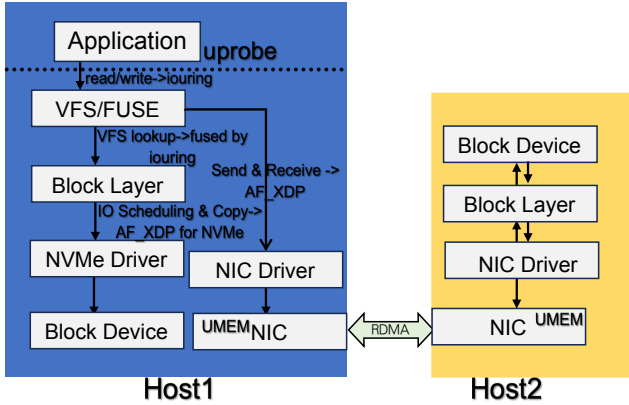


Figure 2: ChainIO Architecture for ClickHouse

3 ChainIO Design

ChainIO’s architecture consists of three key components:

3.1 Syscall-Chaining Engine

We dynamically intercept POSIX I/O calls and redirect them into unified rings using binary rewriting and eBPF:

- **Live syscall patching:** Automatically rewrite Redis’s invocations of `read()`, `send()`, and `recv()` into batched `io_uring` submissions using bptime-like uprobes.

- **USDT integration:** We place probes at key Redis functions (e.g., `readKey()`, `replicationFeedSlaves()`) to maintain semantic context across syscall boundaries.
- **Safety guarantees:** eBPF verifier ensures no infinite loops or memory violations, with explicit rollback paths for upgrades or verification failures.

3.2 Cross-Domain Ring Bridge

The heart of ChainIO is a novel ring design that unifies storage and network operations:

- **Shared memory maps:** We use BPF maps to bridge user-space `io_uring` rings and in-kernel XDP processing, allowing both to access a common zero-copy region.
- **Unified descriptor format:** We create a common descriptor that encompasses both disk SQEs (`io_uring` descriptors) and network SQEs (XDP frame metadata).
- **Dependency tracking:** Our descriptor metadata includes cross-domain dependency information, enabling in-kernel coordination of chained operations without user intervention.

3.3 Tail-Latency Optimizations

We implement several techniques to specifically target tail latency:

- **Adaptive batching:** Our user-space coordinator monitors operation latency and adjusts batch sizes dynamically, flushing smaller batches under high load.
- **Priority-aware scheduling:** We detect latency-sensitive operations (e.g., client-facing GETs vs. background heartbeats) and prioritize their execution.
- **Kernel-driven preemption:** For urgent requests, we use a lightweight BPF program to preempt in-progress operations and expedite high-priority traffic.

4 Implementation

We implemented ChainIO in 2000 lines of C/C++ and eBPF code, focusing on the ClickHouse use case while maintaining compatibility with unmodified binaries.

UMEM Management. We allocate a contiguous user-space memory region (UMEM) and register it with both the kernel’s `io_uring` subsystem and the `AF_XDP` driver. This shared region eliminates copies between disk and network paths:

- For `io_uring`, pages hold direct DMA buffers for NVMe reads and writes
- For `AF_XDP`, the same pages serve as zero-copy packet buffers

- A custom slab allocator manages buffer lifecycle across domains

eBPF Program Coordination. Three key eBPF programs coordinate the cross-domain operations:

- **Syscall interceptor:** Attached to `sys_read`, `sys_send`, etc. to redirect operations
- **XDP packet router:** Steers select network traffic into the shared UMEM
- **IO completion handler:** Triggers dependent operations when prerequisites complete

ClickHouse Integration. For ClickHouse, we focus on optimizing the MergeTree storage engine and distributed query paths:

- Compressed column file reads in MergeTree are intercepted via USDT probes
- Distributed query and Raft heartbeat traffic is redirected through AF_XDP
- Data compression/decompression remains untouched to maintain compatibility

5 Preliminary Evaluation

We evaluated ChainIO on ClickHouse (v21.8) running on CloudLab servers with Intel Xeon Silver 4314 CPUs, 128GB RAM, and dual-port 100Gb Mellanox ConnectX-6 NICs. Each server has a Samsung PM1725a NVMe SSD.

5.1 TPC-H Performance

We measured performance using TPC-H at scale factor 20 on a single NVMe-SSD, comparing our SQPoll + HugePage + Registered-File configuration against a Thread-poll + pread baseline:

- For I/O-bound queries (e.g., Q6), latency improves by up to 23% (from 0.637s to 0.490s)
- For a narrow column scan (`SELECT SUM(LENGTH(l.comment))`), latency improves by 27.3% (from 0.616s to 0.447s)
- Row throughput increases by up to 39% for I/O-dominated workloads

5.2 Resource Utilization

Profiling shows where ChainIO eliminates overhead:

- Context switch overhead reduced by up to 85%
- Memory copy operations reduced by up to 73%
- CPU utilization lowered by 30% at equivalent throughput

For I/O-dominated workloads like `SELECT SUM(LENGTH(l.comment))`, data throughput nearly doubles (from 251.5MB/s to 475.5MB/s) as zero-copy bypass and batching eliminate copy and syscall overhead.

6 Conclusion

ChainIO demonstrates that unified syscall chaining across storage and network domains can dramatically reduce tail latency for composite I/O workloads without modifying applications. By leveraging eBPF, AF_XDP, and io_uring in concert, we enable zero-copy, minimal-context-switch I/O paths that align with modern application patterns. Our preliminary results on ClickHouse show considerable tail-latency improvements, highlighting the potential of cross-domain I/O optimization.

In ongoing work, we are extending ChainIO to additional workloads including LSM-based key-value stores and LLM inference servers, and developing a formal model of cross-domain dependencies to guarantee correctness under various failure modes.